



UNIVERSITÀ DEGLI STUDI DI TRENTO

---

Dipartimento di Ingegneria  
e Scienza dell'Informazione

# Distributed Systems 1

**Gian Pietro Picco**

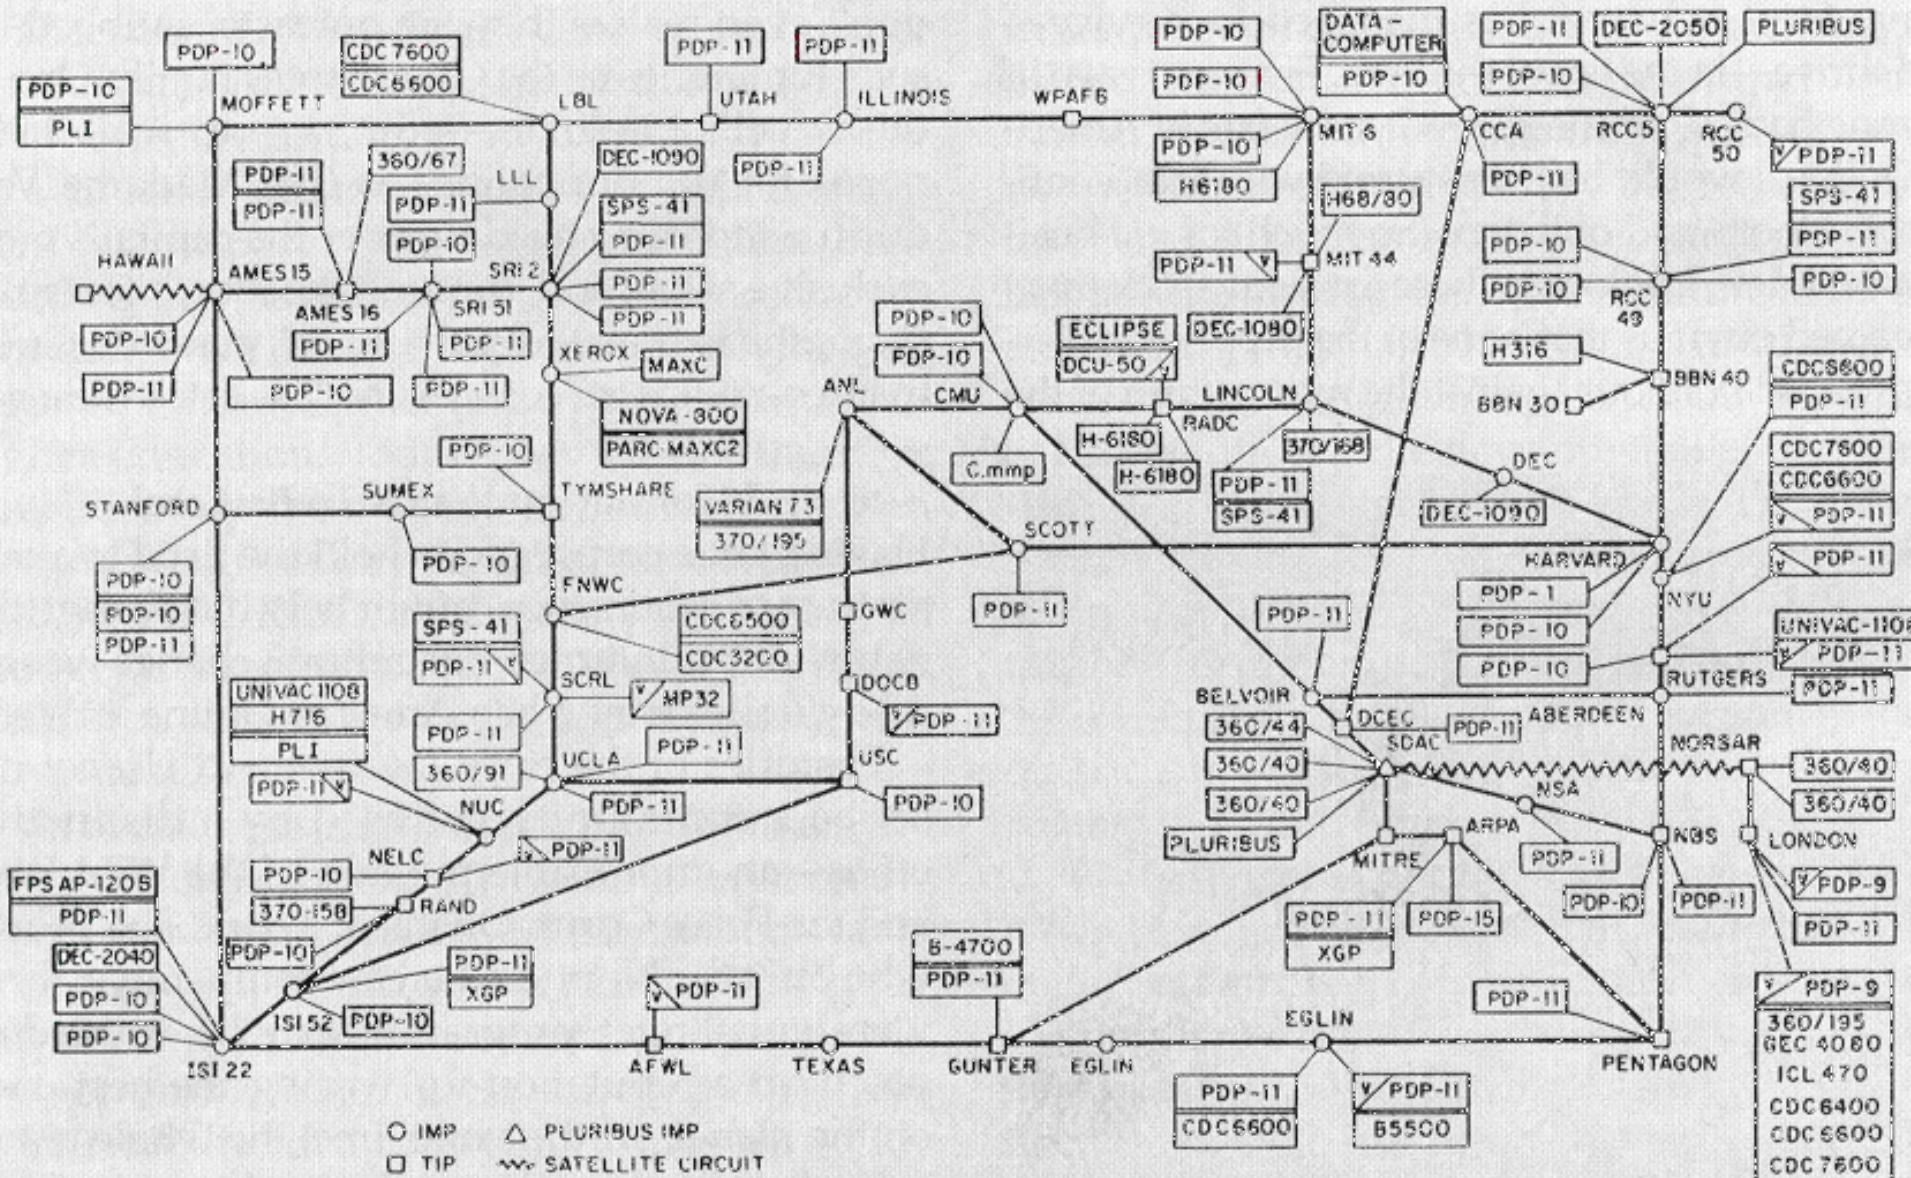
Dipartimento di Ingegneria e Scienza dell'Informazione

University of Trento, Italy

[gianpietro.picco@unitn.it](mailto:gianpietro.picco@unitn.it)

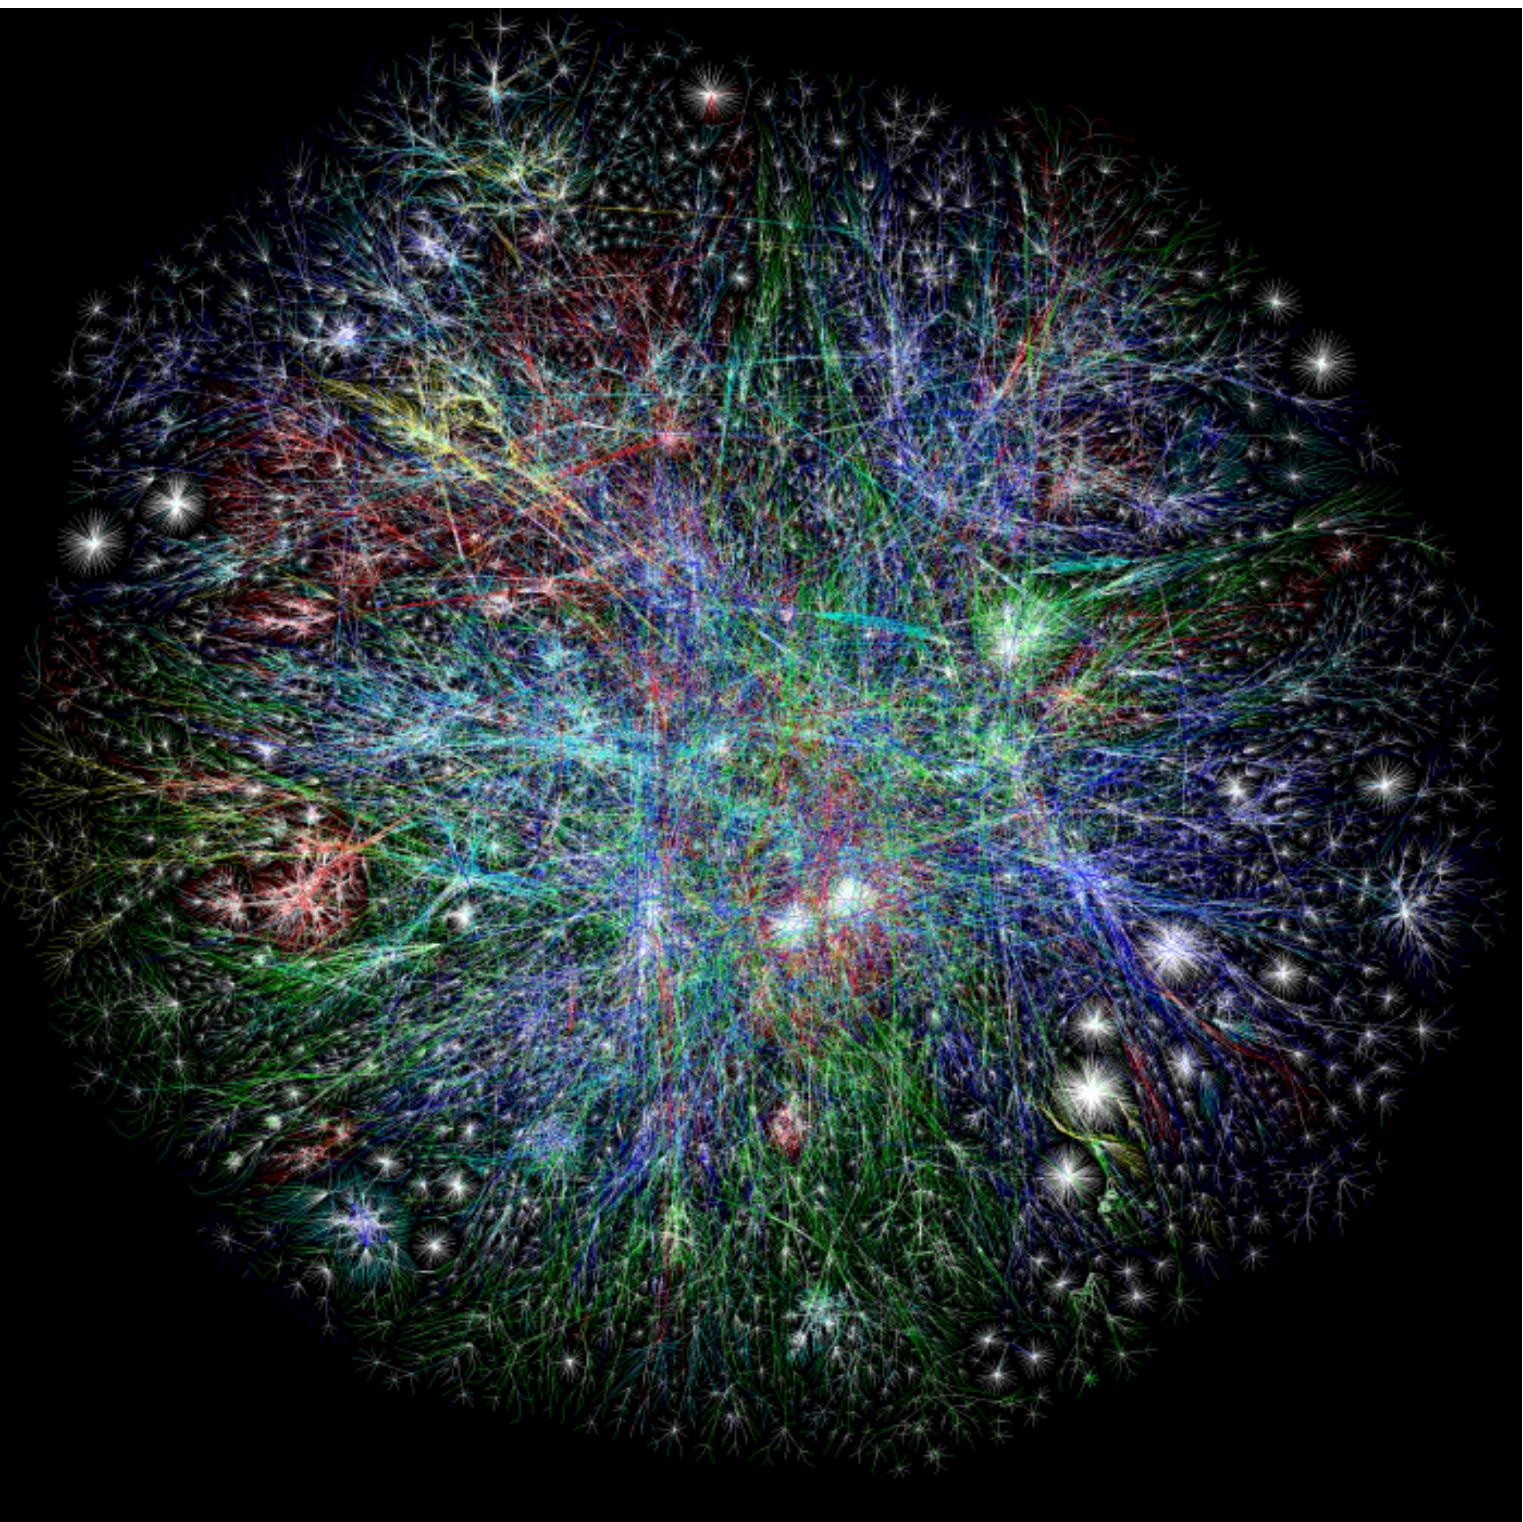
<http://disi.unitn.it/~picco>

**ARPANET LOGICAL MAP, MARCH 1977**

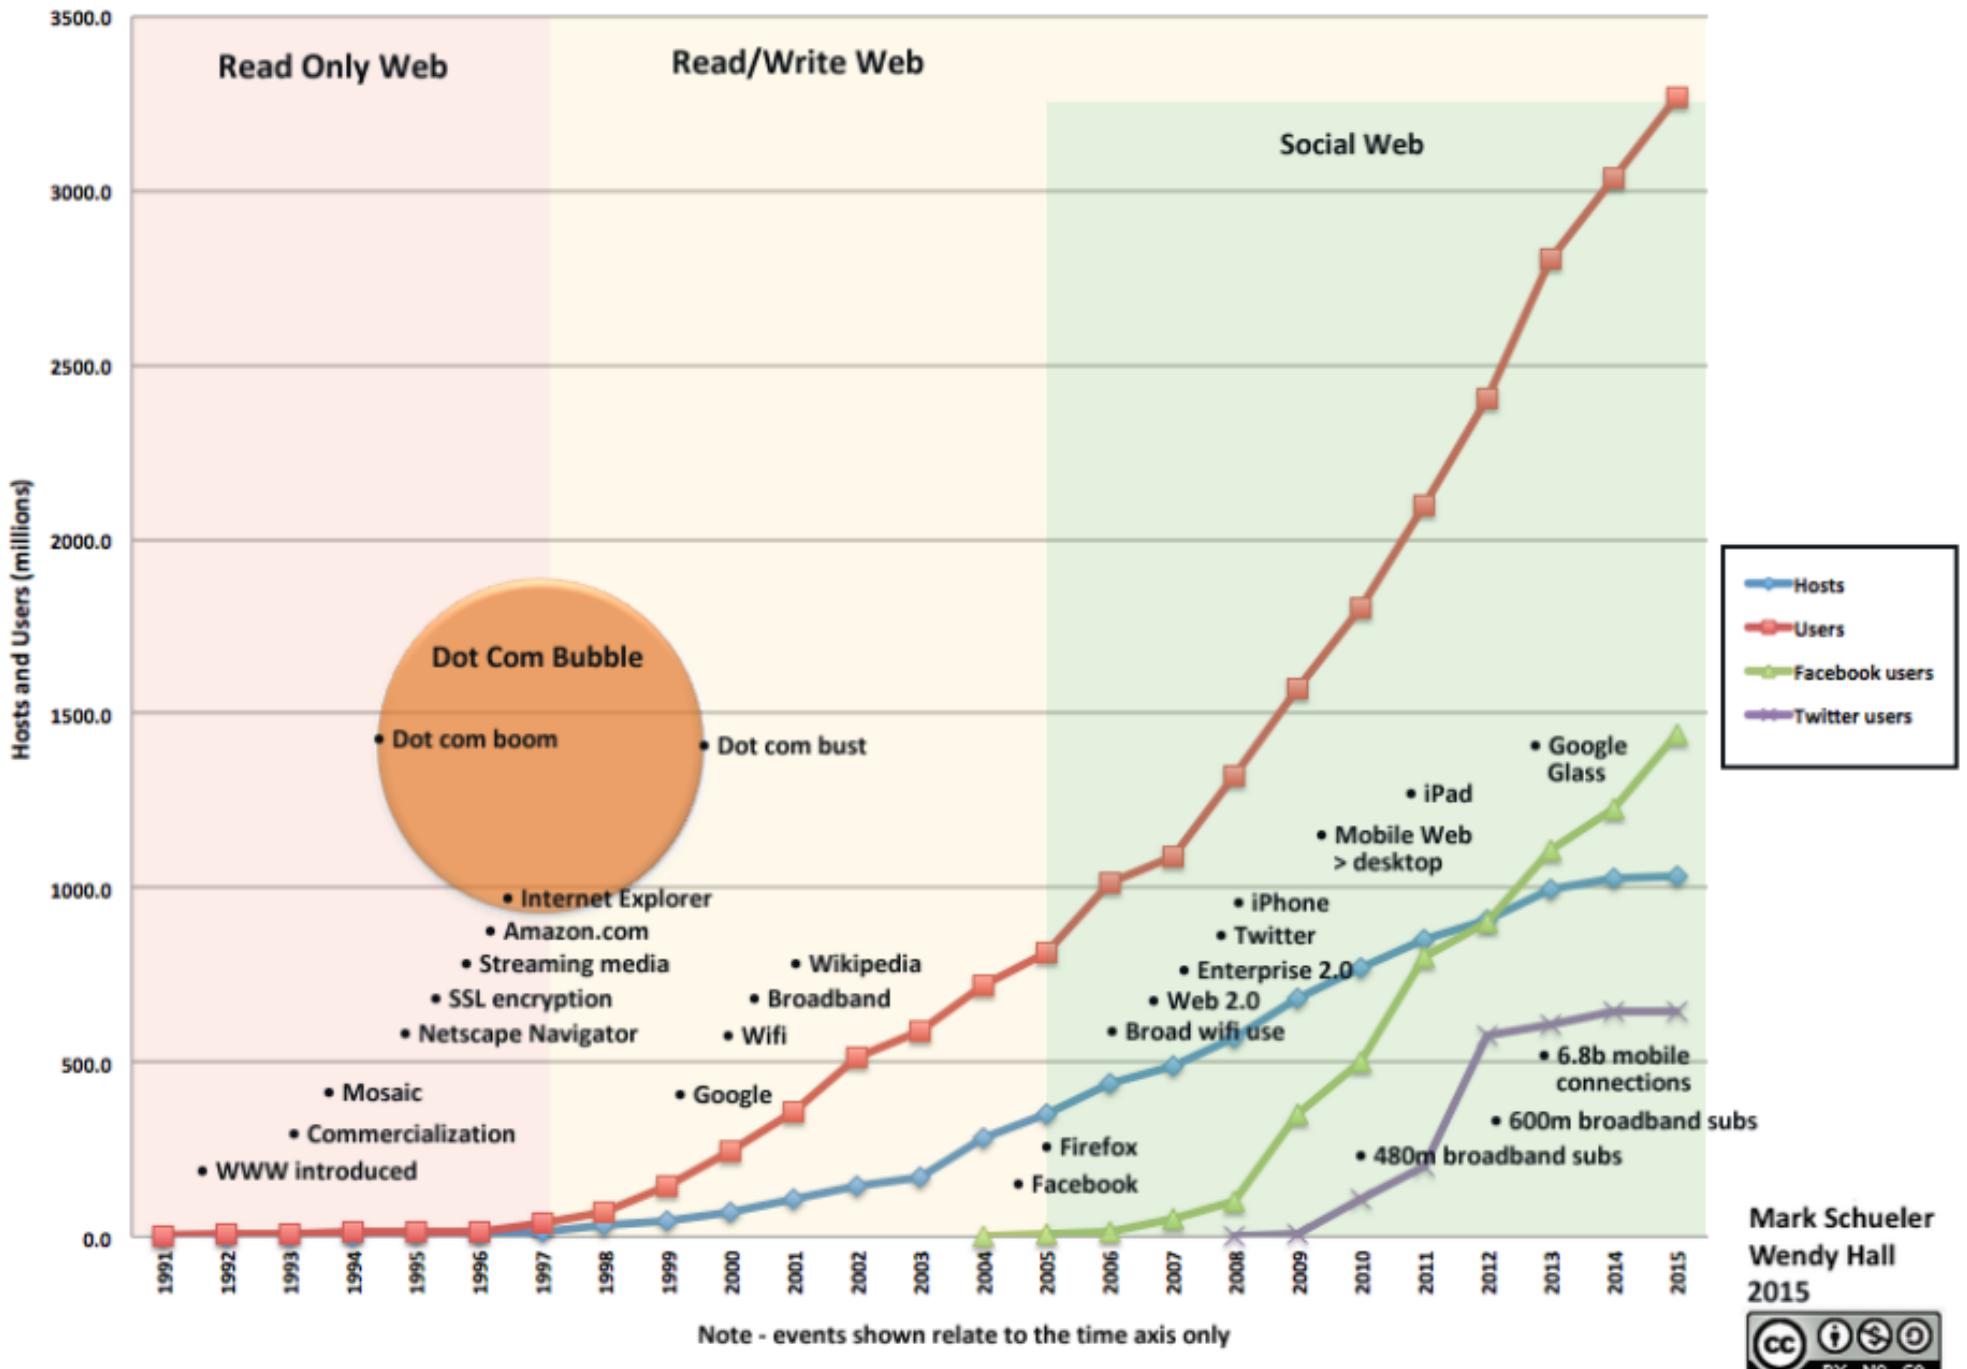


(PLEASE NOTE THAT WHILE THIS MAP SHOWS THE HOST POPULATION OF THE NETWORK ACCORDING TO THE BEST INFORMATION OBTAINABLE, NO CLAIM CAN BE MADE FOR ITS ACCURACY.)

NAMES SHOWN ARE IMP NAMES, NOT (NECESSARILY) HOST NAMES



# Internet Growth - Usage Phases - Tech Events



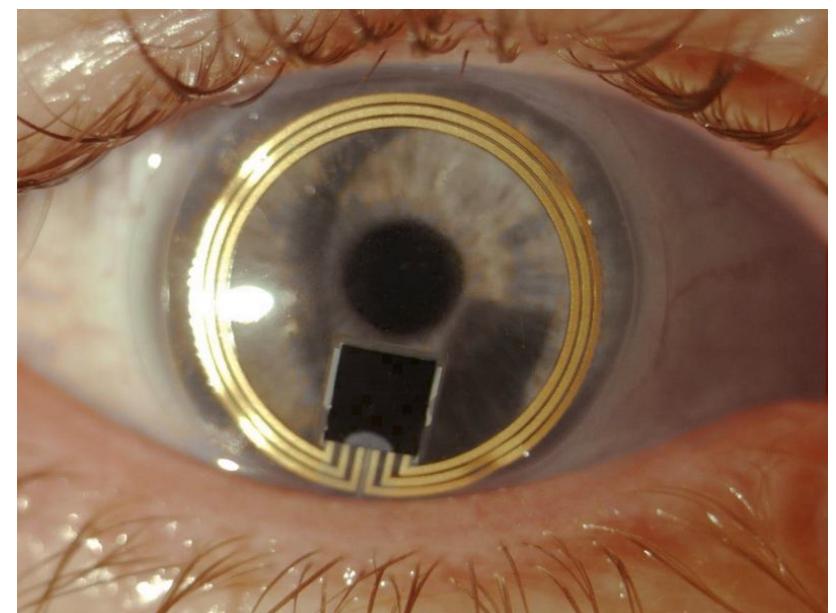
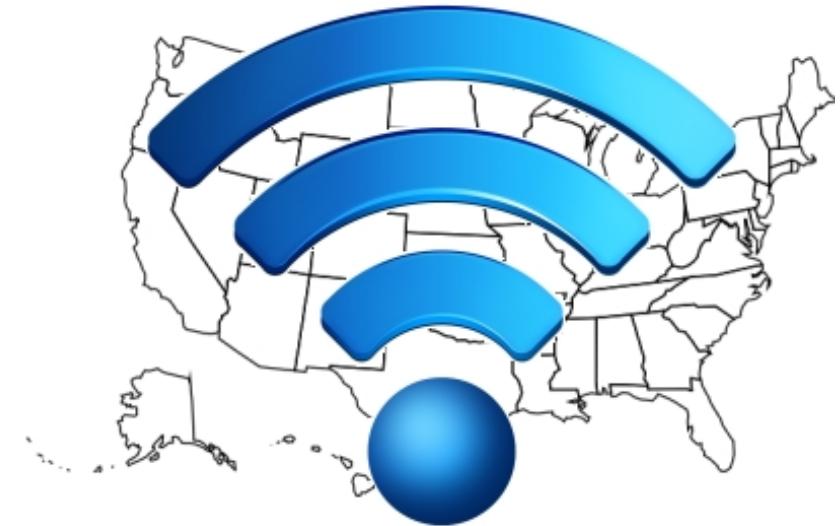
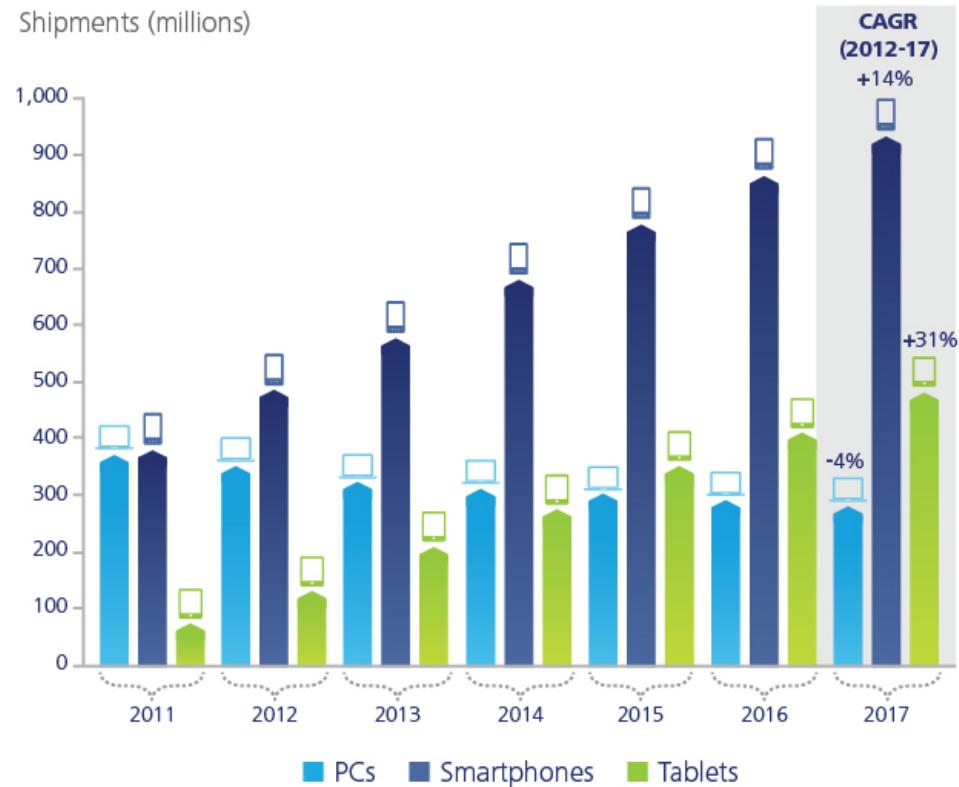
Mark Schueler  
Wendy Hall  
2015



# 2017 This Is What Happens In An Internet Minute



# Recent Game Changers



# Some of Today's Trends

## Data Centers

- Servers are more and more powerful (and cheap)
- They can be “aggregated” to yield high-performance, reliable, scalable computing
  - at a fraction of the cost of a mainframe
  - fundamental role of software and algorithms!
- Data centers are typically composed of thousands machines
  - hosting several virtual machines
  - organized in racks along with high-performance disks
  - connected by high-speed networks
  - with dedicated network stacks

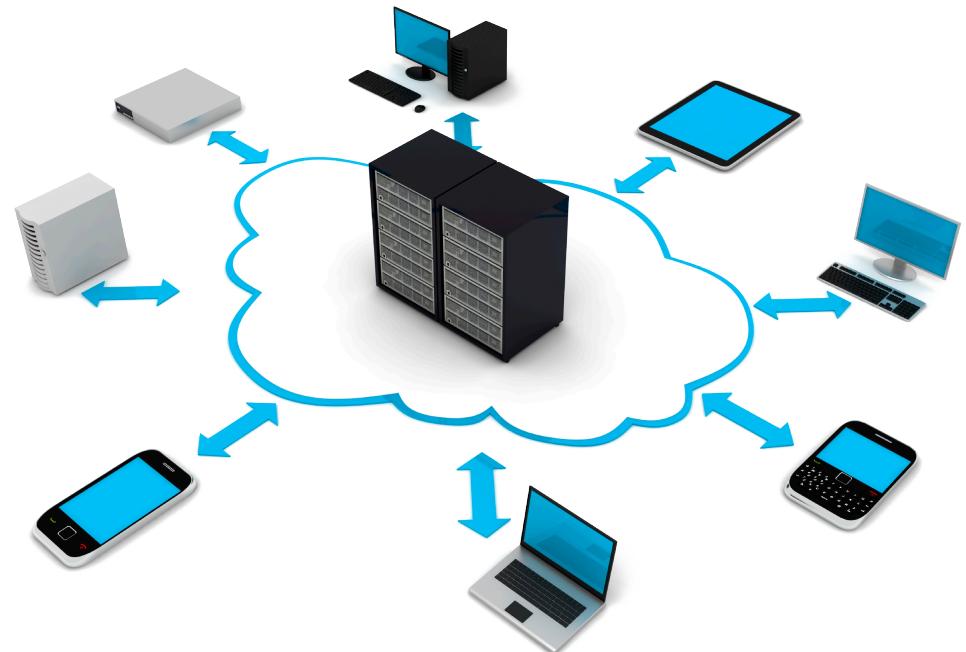
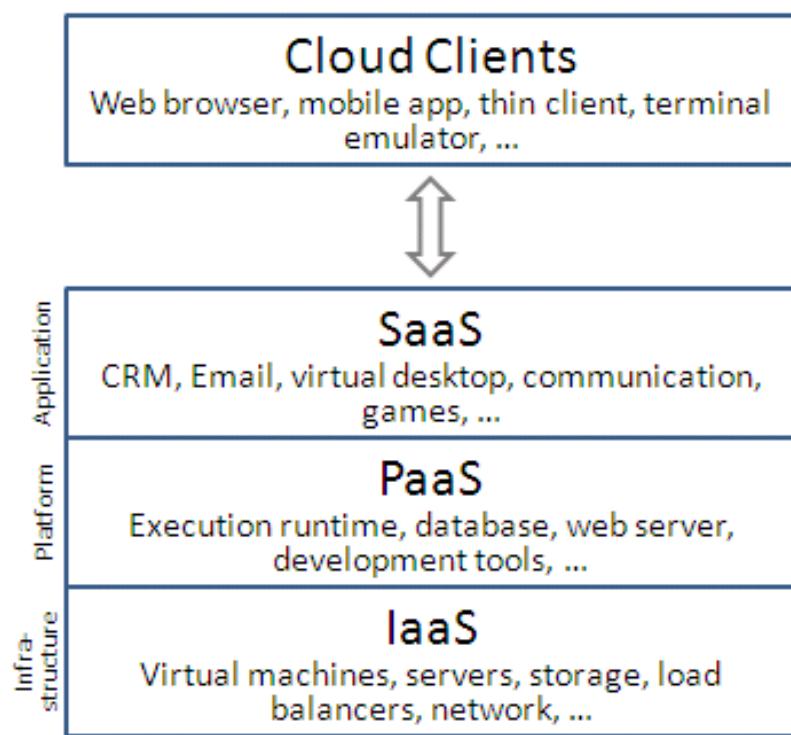
*A glimpse of the UNITN  
data center in Povo2*



# Some of Today's Trends

## Cloud Computing

- **Cloud computing** pushes the concept further, providing (via data centers) users with generic services, at various layers (utility computing)
  - maintenance and operation costs are sustained by the cloud provider instead of the end user



# Some of Today's Trends: *Wearable Computing, Robotics*

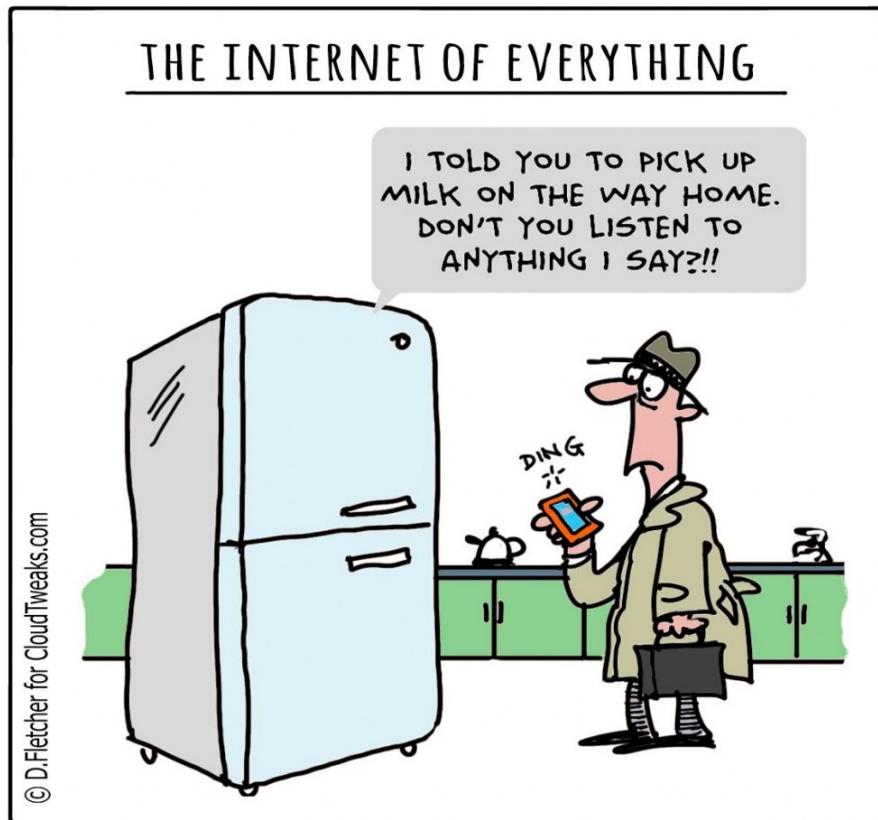


# Some of Today Trends: *Mobile Sensing + Mining Mobile Big Data*



# Some of Today Trends: *The Internet of Things*

- In 2010, the Ericsson CEO publicly stated that, by 2020, there would 50 billion connected devices\*



\* in 2015, they revised the estimate to 28B by 2021

# Libelium Smart World

## Air Pollution

Control of CO<sub>2</sub>, emissions of factories, pollution emitted by cars and toxic gases generated in farms.

## Forest Fire Detection

Monitoring of combustion gases and preemptive fire conditions to define alert zones.

## Wine Quality Enhancing

Monitoring soil moisture and trunk diameter in vineyards to control the amount of sugar in grapes and grapevine health.

## Offspring Care

Control of growing conditions of the offspring in animal farms to ensure its survival and health.

## Sportsmen Care

Vital signs monitoring in high performance centers and fields.

## Structural Health

Monitoring of vibrations and material conditions in buildings, bridges and historical monuments.

## Quality of Shipment Conditions

Monitoring of vibrations, strokes, container openings or cold chain maintenance for insurance purposes.

## Smartphones Detection

Detect iPhone and Android devices and in general any device which works with WiFi or Bluetooth interfaces.

## Perimeter Access Control

Access control to restricted areas and detection of people in non-authorized areas.

## Radiation Levels

Distributed measurement of radiation levels in nuclear power stations surroundings to generate leakage alerts.

## Electromagnetic Levels

Measurement of the energy radiated by cell stations and WiFi routers.

## Traffic Congestion

Monitoring of vehicles and pedestrian affluence to optimize driving and walking routes.



# What Is a Distributed System?

**“pragmatic”**

*A collection of independent, autonomous hosts connected through a communication network.*

*Hosts communicate via message passing to achieve some form of cooperation.*

**“by negation”**

*A parallel system in which there are no:*

- *shared, global clock*
- *shared memory*
- *accurate failure detection*

*“A collection of independent computers that appears to its users as a single coherent system”  
(A.S. Tanenbaum, M. van Steen)*

*“One in which I cannot get any work done because some machine I have never heard of has crashed” (L. Lamport)*

**“optimistic”**

**“pessimistic”**

# Why Distributed Systems?

- Distribution as a *requirement*
  - Applications that require sharing of resources or dissemination of information among geographically distant entities are “natural” distributed systems
  - Examples: banks with several branches, distributed file systems, pervasive systems, ...
- Distribution as a *design/engineering choice*
  - To satisfy non-functional requirements like fault-tolerance, performance vs. cost ratio, quality of service, ...
  - Examples: replicated servers

# Some Defining Features

- ***Concurrency***
  - In centralized systems, concurrency is a design choice
  - In distributed systems, it is a fact of life:  
computers are always (co)-operating at the same time
- ***Absence of a global clock***
  - In centralized systems, the computer's physical clock can be used for the purpose of synchronization
  - In distributed systems, clocks are many and not necessarily synchronized
- ***Absence of a shared memory***
  - In centralized systems, process may coordinate by using a memory area (e.g., a variable) accessible to all of them
  - In distributed systems, the processes reside on different machines with distinct memory address spaces; (asynchronous) message passing is the only means of coordination
- ***Independent (and partial) failures***
  - Centralized systems typically fail completely
  - Distributed systems fail only partially, often due to communication; recovery is complicated by the fact that the application state is itself distributed

# Delays in Distributed Systems

Operation	Delay
CPU calculation	0.000001 ms
Parallel computing (MPI) latency	0.001-0.01 ms
Overhead of method call	0.001-0.01 ms
Wake up thread or process	1 ms
Internet delays	10-1000 ms

- This is one of the reasons why the absence of shared memory introduces a great deal of complexity in distributed systems...

# More on (Distributed) Failures

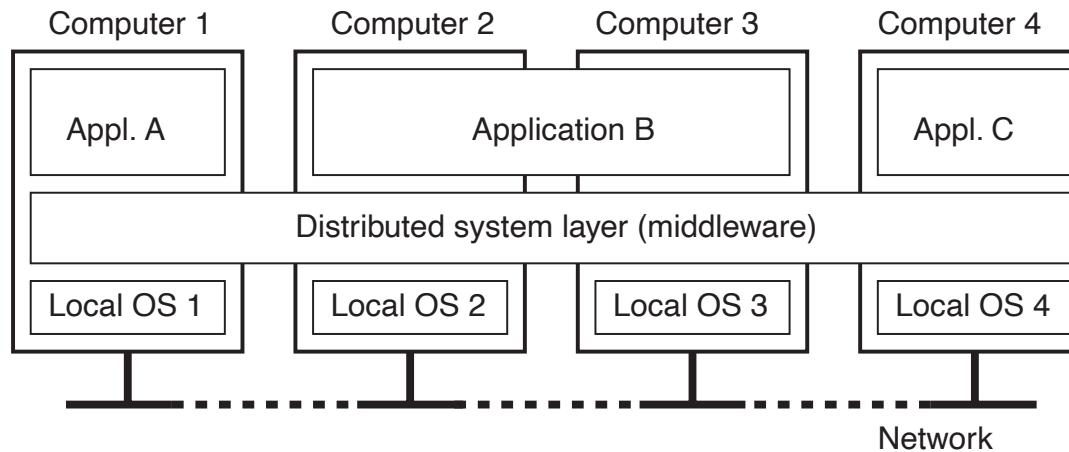
- Failures are a ***problem*** for distributed systems
  - Dependencies among hosts amplify the effect of failures
  - Distributed systems are more complex to design
  - Example:  $n$  dependent hosts, each with probability of failure  $p \rightarrow$  availability is  $(1-p)^n$
- Distributed systems can be a ***solution*** to mitigate failures
  - Providing replicated services with multiple independent processes enables fault tolerance
  - Distributed systems can be made very reliable
  - Example:  $n$  independent hosts, each with probability of failure  $p \rightarrow$  availability is  $1-p^n$

# Parallel vs. Distributed

- **Parallel** systems are composed of multiple processors that:
  - **share memory**
    - used by processors to coordinate their activities
  - **share time**
    - processors access a common clock
  - **share fate**
    - no independent failure
- Massively parallel systems are typically realized via custom hardware architectures
  - Multi-core introduce some amount of parallelism
- The different characteristics of parallel and distributed systems induce different design goals
  - Parallel systems exploit determinism to achieve efficiency
  - Distributed systems are inherently more “loosely coupled” and address the uncertainty this induces

# Challenge: Heterogeneity

- Parallel systems are homogeneous; distributed systems are not, due to differences in
  - network
  - computing hardware
  - operating system
  - programming language
  - multiple implementations (versions)



**Middleware** is a software layer that abstracts from the differences above, providing a uniform computational model

# Goal: Openness

- An open distributed system offers services according to standard rules that describe the syntax and semantics of those services
- These access rules are formalized in (standardized and/or public) protocols or interfaces
- Openness fosters
  - interoperability
    - think about TCP/IP and what it did for the Internet
  - portability
  - extensibility
- Examples: CORBA, JavaEE, Web services, ...

# Challenge & Goal: Security

- Many of the advantages of distributed systems backfire when security is at stake
  - e.g., open standards make it easy for (malicious) nodes to join the distributed system
- Key security aspects:
  - ***confidentiality***: information should be disclosed only to authorized parties
  - ***integrity***: information should be altered only by authorized parties; unauthorized alterations should be detectable and recoverable
  - ***availability***: services should be provided in spite of malicious behavior
- We will not discuss further these aspects; see other courses in the curriculum

# Challenge: Failure Handling

- Dealing with failures is key in distributed systems
- Several techniques exist, with different goals:
  - failure detection (e.g., message checksums)
  - failure masking (e.g., email retransmissions)
  - failure tolerance (e.g., replicated servers)
  - failure recovery (e.g., log files)
- They are not always applicable (e.g., masking)
- Real systems use a combination of the above, depending on the application requirements and fault model

# Goal: Scalability

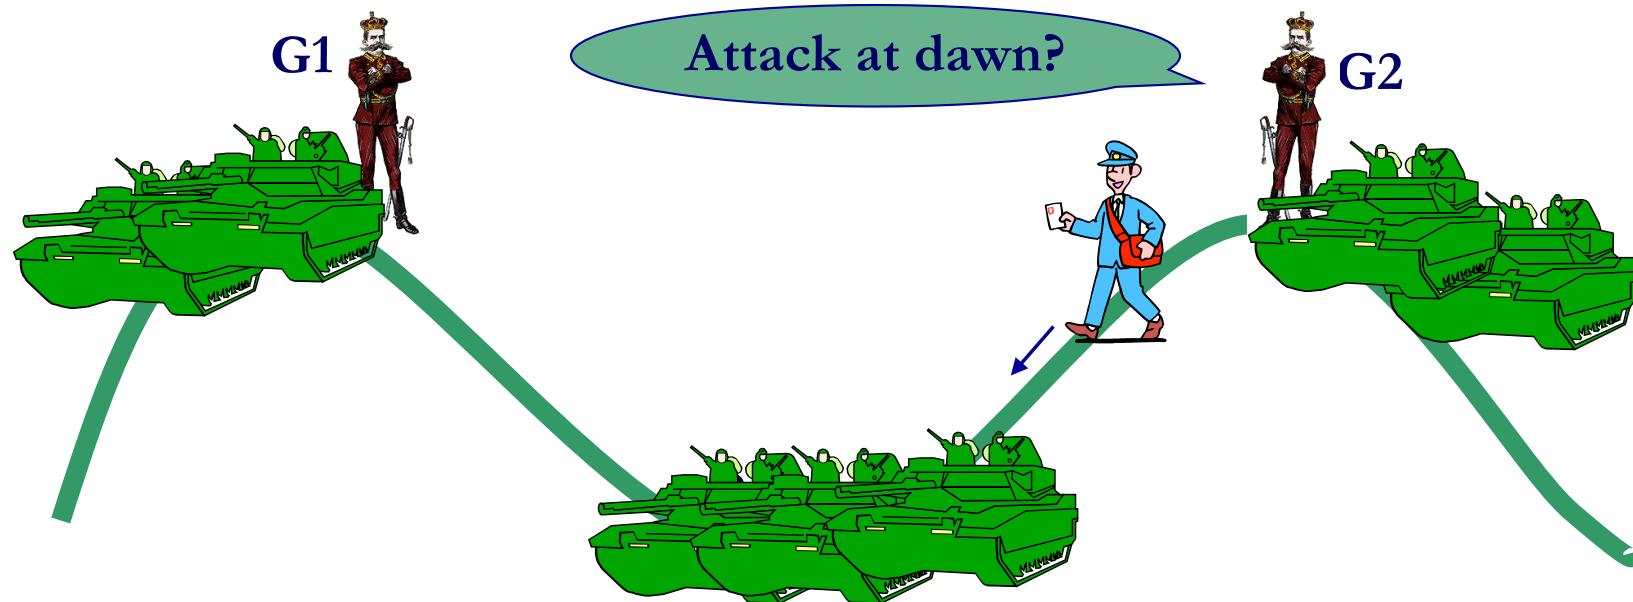
- It is the ability to increase the scale of the system in terms of:
  - size, i.e., number of users, hosts, or resources
  - geographic span
  - administrative span
- Tradeoff between scale and performance
- Centralized solutions are easier to design and reason about, but may yield to
  - a performance bottleneck...
  - ... and a single point of failure
  - centralization of services, data, algorithms
- Decentralized solutions are desirable
  - e.g., using asynchronous communications, hierarchical structuring, caching and replication, ...
- In practice, systems often use a mix
  - true for e.g., several services in the Google back-end

# Goal: Transparency

Transparency	Description
<b>Access</b>	Hide differences in data representation and how a resource is accessed
<b>Location</b>	Hide where a resource is located
<b>Migration</b>	Hide that a resource may move to another location
<b>Replication</b>	Hide that a resource may exist in multiple copies
<b>Concurrency</b>	Hide that a resource may be shared by several competitive users
<b>Failure</b>	Hide the failure and recovery of a resource

- Some things cannot be made transparent
  - Timezones, communication delays
- Hiding too much may have a strong, negative performance impact
  - E.g., accessing repeatedly a remote object without knowing

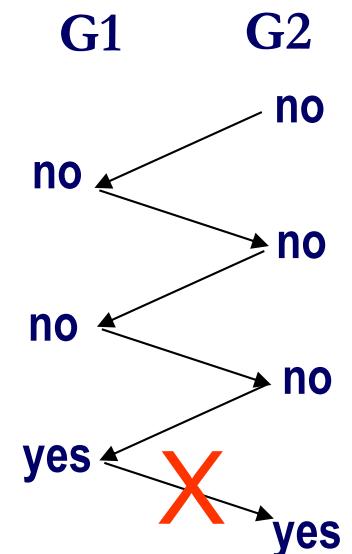
# A Seemingly Easy Problem ...



- Generals decide separately whether or not to attack
  - If **both** attack: success
  - If only one attacks: failure
- Messengers are employed for coordinating the decision
  - But they can be lost/captured/delayed
- **Question:** How do the generals ensure that they take the same decision? What algorithm do they use?

# (Theoretical) Impossibility of Distributed Consensus

- It is impossible to design an algorithm to guarantee success or no attack if the messengers can be lost:
  - **Impossibility of distributed consensus**
  - Formally demonstrated (Fischer, Lynch, Patterson, 1985)
- (Informal) proof:
  - Reaching an agreement on one of two possible decisions requires the successful arrival of at least one message
  - Consider scenario A in which the **fewest** delivered messages that will result in agreement to attack are delivered
  - Let scenario B be the same as A except that the last message delivered in A is lost in B, and any other messages that might be sent later are also lost
  - Suppose this last message is from General 1 to General 2
  - General 1 sees the same messages in both scenarios, so he definitely attacks
  - However, the **minimality** assumption of A implies that General 2 cannot also decide to attack in scenario B, so he must make a different decision
  - Hence the problem is unsolvable



*What  
does it mean  
in practice?*

# Distributed Consensus in Practice

- Does it really matter in real life? **Yes!!!**
  - E.g., when you withdraw 100 EUR (at the ATM), your balance (at the bank) should be decreased by 100 EUR
    - This is the “commit” of a transaction in a distributed database, over a communication network
  - Agree on values of replicated, distributed sensors
  - Agree on whether a system component is faulty
- Then we cannot design distributed systems? **No!!!**
- In practice, real systems either
  - Change the assumptions
    - E.g., make links reliable (enough)
  - Reduce the guarantees:
    - E.g., only probabilistic instead of deterministic

*In theory, theory and practice are the same.  
In practice, they are not.*  
-- Yogi Berra



# The Course

- Goals:
  - lectures:
    - discuss fundamental concepts in distributed systems
  - labs:
    - understand some of these concepts hands-on ...
    - ... while experiencing with software technology for developing distributed systems
- Macro-topics:
  - synchronization
  - replication and consistency
  - fault tolerance
  - middleware